Try again once you are ready.

Required to pass: 80% or higher

You can retake this quiz up to 3 times every 8 hours.

Back to Week 3

Retake

0 / 1 point

1.

Which of the following are true? (Check all that apply.)

- [ ] a is the activation output of the nd layer for the th training example

Un-selected is correct

- [ ] a denotes activation vector of the th layer on the nd training example.

Un-selected is correct

- [x] a denotes the activation vector of the nd layer.

Correct

- [x] X is a matrix in which each column is one training example.

Correct

- [ ] a denotes the activation vector of the nd layer for the th training example.

This should be selected

- [x] a is the activation output by the th neuron of the nd layer

Correct

- [ ] X is a matrix in which each row is one training example.

Un-selected is correct

---

1 / 1 point

2.
The tanh activation usually works better than sigmoid activation function for hidden units because the mean of its output is closer to zero, and so it centers the data better for the next layer. True/False?

- True

Correct
Yes. As seen in lecture the output of the tanh is between -1 and 1, it thus centers the data which makes the learning simpler for the next layer.

- False

---

1 / 1 point

3.
Which of these is a correct vectorized implementation of forward propagation for layer l, where l L?

- Z^l $W^l A^l$ b^l
- A^l g^l Z^l

- Z^l $W^l A^l$ b^l
- A^l g^l Z^l

- Z^l W^l A^l b^l
- A^l g^l Z^l

- Z^l $W^l A^l$ b^l
- A^l g^l Z^l

Correct

---

0 / 1 point

4.
You are building a binary classifier for recognizing cucumbers (y=1) vs. watermelons (y=0). Which one of these activation functions would you recommend using for the output layer?

ReLU

This should not be selected
No. Sigmoid outputs a value between 0 and 1 which makes it a very good choice for binary classification. You can classify as 0 if the output is less than 0.5 and classify as 1 if the output is more than 0.5. It can be done with tanh as well but it is less convenient as the output is between -1 and 1.

Leaky ReLU

sigmoid

tanh

---

1 / 1 point

5.

Consider the following code:

```
A = np.random.randn(4,3)
B = np.sum(A, axis = 1, keepdims = True)
```

What will be B.shape? (If you're not sure, feel free to run this in python to find out).

- (4, 1)

Correct
Yes, we use (keepdims = True) to make sure that A.shape is (4,1) and not (4,). It makes our code more rigorous.

- (4,)
- (, 3)
- (1, 3)

---

1 / 1 point

6.
Suppose you have built a neural network. You decide to initialize the weights and biases to be zero. Which of the following statements is true?

- Each neuron in the first hidden layer will perform the same computation. So even after multiple iterations of gradient descent each neuron in the layer will be computing the same thing as other neurons.

Correct

- Each neuron in the first hidden layer will perform the same computation in the first iteration. But after one iteration of gradient descent they will learn to compute different things because we have "broken symmetry".

Each neuron in the first hidden layer will compute the same thing, but neurons in different layers will compute different things, thus we have accomplished “symmetry breaking” as described in lecture.

The first hidden layer’s neurons will perform different computations from each other even in the first iteration; their parameters will thus keep evolving in their own way.

✗ 0 / 1 point

7.
Logistic regression’s weights w should be initialized randomly rather than to all zeros, because if you initialize to all zeros, then logistic regression will fail to learn a useful decision boundary because it will fail to “break symmetry”, True/False?

True

This should not be selected
No, Logistic Regression doesn't have a hidden layer. If you initialize the weights to zeros, the first example x fed in the logistic regression will output zero but the derivatives of the Logistic Regression depend on the input x (because there's no hidden layer) which is not zero. So at the second iteration, the weights values follow x's distribution and are different from each other if x is not a constant vector.

False

✓ 1 / 1 point

8.
You have built a network using the tanh activation for all the hidden units. You initialize the weights to relative large values, using np.random.randn(..,..)*1000. What will happen?

This will cause the inputs of the tanh to also be very large, thus causing gradients to also become large. You therefore have to set α to be very small to prevent divergence; this will slow down learning.

This will cause the inputs of the tanh to also be very large, causing the units to be “highly activated” and thus speed up learning compared to if the weights had to start from small values.

It doesn't matter. So long as you initialize the weights randomly gradient descent is not affected by whether the weights are large or small.

This will cause the inputs of the tanh to also be very large, thus causing gradients to be close to zero. The optimization algorithm will thus become slow.

Correct
Yes. tanh becomes flat for large values, this leads its gradient to be close to zero. This slows down the optimization algorithm.

---

0 / 1 point

9.
Consider the following 1 hidden layer neural network:

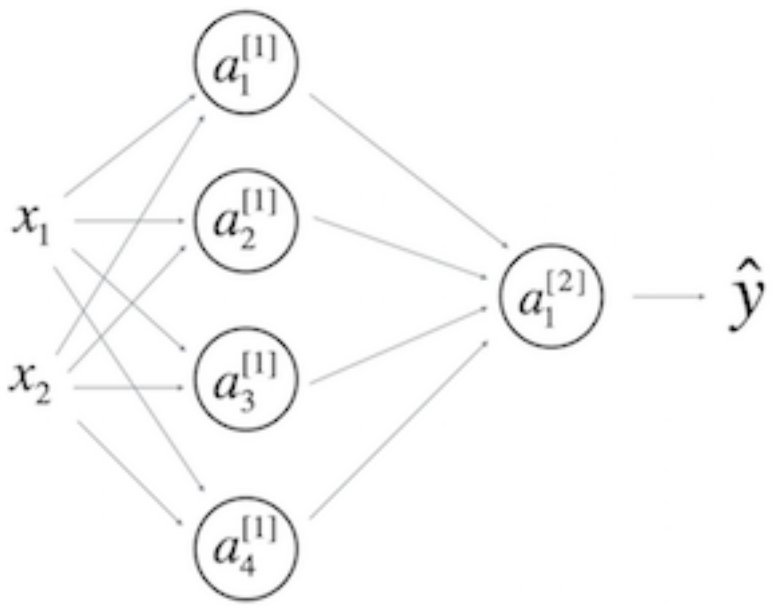

Which of the following statements are True? (Check all that apply).

- [x] W will have shape (2, 4)

This should not be selected

- [ ] b will have shape (4, 1)

This should be selected

- [ ] W will have shape (4, 2)

This should be selected

- [ ] b will have shape (2, 1)

This should not be selected

- [ ] W will have shape (1, 4)

This should be selected

- [ ] b will have shape (4, 1)

Un-selected is correct

- [ ] W will have shape (4, 1)

Un-selected is correct

- [ ] b will have shape (1, 1)

This should be selected

---

0 / 1 point

10.
In the same network as the previous question, what are the dimensions of Z and A ?

- Z and A are (4,2)
- Z and A are (1,4)
- Z and A are (4,1)
- Z and A are (4,m)